

ECON 626 Prediction Competition 6 Code

Objective: To predict using classification which of the 15,578 text snippets from news articles include the words "economic","economy", "economics" or the character sequence "econom" (not case-sensitive) in it.

Importing libraries

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")

from sklearn.preprocessing import LabelEncoder

import nltk
from nltk.stem import PorterStemmer
from nltk import LancasterStemmer
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.metrics import edit_distance
import string
import re
from bs4 import BeautifulSoup

from wordcloud import WordCloud

from collections import Counter

from sklearn.model_selection import train_test_split

from sklearn.feature_extraction.text import CountVectorizer

from sklearn.feature_extraction.text import TfidfTransformer

from sklearn import preprocessing

from sklearn import svm
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report

from sklearn.model_selection import GridSearchCV

from scipy.sparse import hstack

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
from tensorflow.keras.utils import to_categorical
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Input
```

Importing data

```
In [ ]: data_path = "/Users/andrew/Downloads/UW courses/ECON 626/Prediction Competit
text_df= pd.read_csv(data_path)
#create a dataframe for our text dataset
```

Note I made two alterations to the dataset pre importing. They are listed below:

- First I added a header to the dataset called 'text'
- After searching for duplicates I found that some of the text started with "--" which created an error with the csv after saving where some of the values turned to '#NAME'
- I deleted a '<' in line 12575 in the data since the symbol was causing an error where it believed the text inside the symbol was html text and missed an 'econom' sequence. I will look into switch from an 'anything in a <>' regex to a 'specific html commands' regex since this could cause problems in predictions.

```
In [ ]: text_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15578 entries, 0 to 15577
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    15578 non-null     object
dtypes: object(1)
memory usage: 121.8+ KB
```

```
In [ ]: text_df.head()
```

```
Out [ ]:
```

	text
0	CEA believe that some special factors, such as...
1	approved a \$282.6 billion measure to fund the ...
2	(AP) ≡ŲŲ An unexpectedly quick rise in. the pr...
3	\$17.28 million would be put into a reserve for...
4	Reserve policy makers voted to keep short-term...

Data Cleaning

```
In [ ]: num_na = text_df.isna().sum()
num_na
```

```
Out[ ]: text    0
        dtype: int64
```

```
In [ ]: duplicate_rows = text_df.duplicated(keep=False)

# Check if there are any duplicates
if duplicate_rows.any():
    print(text_df[duplicate_rows])
    # Print the rows that are duplicates
else:
    print("No duplicates found.")
```

```

text
506  general manager, said Silicon Valley's patent ...
1532 yet to damp innovation in Silicon Valley, at l...
1559 way in what central banks are and what they do...
3469 general manager, said Silicon Valley's patent ...
4255 bond market intensified Thursday, as the gap b...
5837 the 1990s, when government spending and taxes ...
7375 Americans to gamble that the U.S. economy can ...
7539 yet to damp innovation in Silicon Valley, at l...
8196 inflation fears on the rise, investors may wan...
8563 from the recession and the collapsed real esta...
8621 bond market intensified Thursday, as the gap b...
9100 Americans to gamble that the U.S. economy can ...
10721 way in what central banks are and what they do...
11204 on both domestic and foreign policy.</br></br>...
11591 from the recession and the collapsed real esta...
12441 inflation fears on the rise, investors may wan...
14996 on both domestic and foreign policy.</br></br>...
15460 the 1990s, when government spending and taxes ...
```

```
In [ ]: text_df['length'] = text_df['text'].apply(len)
        text_df
```

Out []:

	text	length
0	CEA believe that some special factors, such as...	644
1	approved a \$282.6 billion measure to fund the ...	556
2	(AP) ¶ An unexpectedly quick rise in. the pr...	453
3	\$17.28 million would be put into a reserve for...	469
4	Reserve policy makers voted to keep short-term...	608
...
15573	Grover Norquist says emphatically and repeated...	423
15574	in die economy, Laura D¶Andrea Tyson, chairm...	580
15575	of living are not rising noticeably. Many of t...	833
15576	economy merely flexes some long-unused muscles...	767
15577	Video Investments, an investment-research and ...	824

15578 rows x 2 columns

Data Preprocessing

In []: `sample = text_df['text'].iloc[12573]`
`sample`

Out []: 'a recesbave been denied their con-sion or is heading t0wardtitutional ri h
 ts by such</br></br>Leading indicators -leavåfnmates 1ft ers and the little
 doubt as to which way)Uyn,nfåÊ letters addressed the economy is headed intå
 i inmates.</br></br>1970\x89Û0toward a recession.\x89Û\x9d Also, it is char
 ged that J. Gordon Gifford, editolphotostatic copies of letters of the comm
 ission\x89Ûas monthly inmates have been made report, said all six of th%nd
 turned over to the attor-leading indicators were of|,ey general\x89Ûas offi
 ce, in some in January from Decembeijnstances and below the level of Janu-
 \t..\t\x89Û_ ary, 1969. He said they had/fhe suit alleges prison of-been de
 clining since Septemficials have infringed upon ber.'

In []: *# Test on a sample text*

```
stemmer = PorterStemmer()
lemmatizer = WordNetLemmatizer()

text = text_df['text'].iloc[12573]

# Remove HTML tags using BeautifulSoup
#clean_text = BeautifulSoup(text).get_text()

#clean_text = ' '.join(re.split(r'<.*?>', clean_text))

clean_text = re.sub("<.*?>", lambda m: " " * len(m.group(0)), text)
```

```

clean_text = ''.join([char for char in text if char not in string.punctuation])

# # Replace specific emojis with corresponding text
# clean_text = re.sub(r'<3', '<heart>', clean_text)
# clean_text = re.sub(r"[8:=;]['\-]?[()d]+", '<smile>', clean_text)
# clean_text = re.sub(r"[8:=;]['\-]?[+]", '<sadface>', clean_text)
# clean_text = re.sub(r"[8:=;]['\-]?[\/|l*]", '<neutralface>', clean_text)
# clean_text = re.sub(r"[8:=;]['\-]?[p+]", '<lolface>', clean_text)

# # Remove non-alphabetical symbols
clean_text = re.sub('[^A-Za-z ]+', '', clean_text)

# # Remove stopwords
clean_text = ' '.join([word.lower() for word in clean_text.split() if word.lower() not in stopwords])
clean_text = [stemmer.stem(lemmatizer.lemmatize(word.lower())) for word in clean_text]

print(clean_text)

```

```

['recesbav', 'deni', 'consion', 'head', 'twardtitut', 'ri', 'ht', 'suchbrbrl',
'ead', 'indic', 'leavnmat', 'ft', 'er', 'littl', 'doubt', 'wayuynnf', 'lette',
'r', 'address', 'economy', 'head', 'int', 'inmatesbrbrtoward', 'recess', 'also',
'o', 'charg', 'j', 'gordon', 'gifford', 'editolphotostat', 'copi', 'letter',
'commiss', 'monthli', 'inmat', 'made', 'report', 'said', 'six', 'thnd', 'turn',
'n', 'attorlead', 'indic', 'ofey', 'gener', 'offic', 'januari', 'decembeijns',
't', 'level', 'janu', 'ari', 'said', 'hadfh', 'suit', 'alleg', 'prison', 'ofb',
'een', 'declin', 'sinc', 'septemfici', 'infring', 'upon', 'ber']

```

```

In [ ]: # Create a function
def text_processing(text:str) -> list:

    clean_text = re.sub("<.*?>", lambda m: " " * len(m.group(0)), text)
    clean_text = ''.join([char for char in clean_text if char not in string.punctuation])
    clean_text = re.sub("\s+", " ", clean_text)
    clean_text = re.sub("<.*?>", ' ', clean_text)
    clean_text = re.sub(r'[^w ]+', '', clean_text)

    # Convert the text to lowercase for case-insensitive matching
    clean_text_lower = clean_text.lower()

    # Remove non-alphabetical symbols
    clean_text = re.sub('[^A-Za-z ]+', '', clean_text_lower)

    # Remove stopwords
    clean_text = ' '.join([word for word in clean_text.split() if word.lower() not in stopwords])

    # Stem and lemmatize the tokens
    stemmer = LancasterStemmer()
    lemmatizer = WordNetLemmatizer()
    clean_text_tokens = [lemmatizer.lemmatize(stemmer.stem(word)) for word in clean_text.split()]

    # Define the words and character sequence to check for
    words_to_check = ["economic", "economy", "economics", "econom", "economy"]

    # Check if any of the words or character sequence are present in the tokens

```

```
for word in clean_text_tokens:
    for word_to_check in words_to_check:
        if edit_distance(word, word_to_check) <= 2:
            return clean_text_tokens # Return the corrected tokens if t

return clean_text_tokens # Return the original tokens if none of the wc
```

```
In [ ]: text_processing(text_df['text'].iloc[5500])
```

```
Out[ ]: ['expect',
        'giv',
        'way',
        'renew',
        'adv',
        'short',
        'turnaround',
        'com',
        'liv',
        'cost',
        'resum',
        'upward',
        'march',
        'exceiv',
        'new',
        'car',
        'pric',
        'exempl',
        'lik',
        'edg',
        'model',
        'rol',
        'produc',
        'lin',
        'deal',
        'custom',
        'shav',
        'pric',
        'model',
        'year',
        'progress',
        'year',
        'poor',
        'sal',
        'perform',
        'enco',
        'many',
        'deal',
        'giv',
        'ev',
        'big',
        'discount',
        'u',
        'word',
        'detroit',
        'model',
        'cost',
        'year',
        'car',
        'mean',
        'automobl',
        'pric',
        'jump',
        'isharply',
        'next',
        'fal']
```

```
In [ ]: text_df['tokens'] = text_df['text'].apply(text_processing)
```

```
In [ ]: text_df.head()
```

```
Out[ ]:
```

	text	length	tokens
0	CEA believe that some special factors, such as...	644	[cea, believ, spec, fact, hug, runup, stock, m...
1	approved a \$282.6 billion measure to fund the ...	556	[approv, bil, meas, fund, nat, defens, com, fi...
2	(AP) ≡ŮÓ An unexpectedly quick rise in the pr...	453	[ap, unexpect, quick, ri, prim, rat, many, nat...
3	\$17.28 million would be put into a reserve for...	469	[mil, would, put, reserv, next, year, ev, mone...
4	Reserve policy makers voted to keep short-term...	608	[reserv, policy, mak, vot, keep, shortterm, in...

```
In [ ]: # Define a function to check if the specified words or character sequence are
def check_economic_presence(tokens):
    # Define the words and character sequence to check for
    words_to_check = ["economic", "economy", "economics", "econom", "economy"]

    # Define the regular expression pattern to search for
    pattern = re.compile(r'econom', re.IGNORECASE)

    # Check if any of the words or character sequence are present in the tokens
    for word in tokens:
        if word in words_to_check or pattern.search(word):
            return 1 # Return 1 if any of the words or character sequence are found

    return 0 # Return 0 if none of the words or character sequence are found

# Apply the function to the 'tokens' column to create the 'class' column
text_df['class'] = text_df['tokens'].apply(check_economic_presence)
```

```
In [ ]: text_df.head(25)
```


Out []:

	text	length	tokens	class
0	CEA believe that some special factors, such as...	644	[cea, believ, spec, fact, hug, runup, stock, m...	1
1	approved a \$282.6 billion measure to fund the ...	556	[approv, bil, meas, fund, nat, defens, com, fi...	0
2	(AP) ≡ŮŮ An unexpectedly quick rise in. the pr...	453	[ap, unexpect, quick, ri, prim, rat, many, nat...	0
3	\$17.28 million would be put into a reserve for...	469	[mil, would, put, reserv, next, year, ev, mone...	0
4	Reserve policy makers voted to keep short-term...	608	[reserv, policy, mak, vot, keep, shortterm, in...	0
5	the lender asked us if we wanted to lock in.</...	824	[lend, ask, u, want, lock, sint, first, hous, ...	0
6	called off its effort to drive down the value ...	752	[cal, effort, driv, valu, doll, cur, return, a...	1
7	striking unions at. the Greenbrier Hotel resor...	421	[striking, un, greenbry, hotel, resort, whit, ...	0
8	short-term political pressure. But an agency t...	596	[shortterm, polit, press, ag, alloc, credit, s...	0
9	typical recession would drop enough to make st...	797	[typ, recess, would, drop, enough, mak, stock,...	0
10	the standard deduction, personal exemption and...	476	[standard, deduc, person, exempt, cutoff, fig,...	0
11	nation's factories continued to fall last mont...	984	[nat, fact, continu, fal, last, mon, provid, s...	1
12	percentage of those entitled to unemployment i...	971	[perc, entitl, unemploy, in, remain, per, cent...	0
13	the aid earlier than planned threatens to dimi...	1038	[aid, ear, plan, threatens, dimin, impact, lar...	0
14	oversaw the conclusion of his final General As...	729	[oversaw, conclud, fin, gen, assembl, sess, la...	1
15	and labor costs rose sharply in the second qua...	740	[lab, cost, ro, sharply, second, quart, spark,...	0
16	But companies still pay, to the tune of \$1.6 b...	540	[company, stil, pay, tun, bil, last, year, pla...	0
17	securities were narrowly mixed in quiet tradin...	618	[sec, narrow, mix, quiet, trad, many, biggest,...	0
18	Japan's real gross domestic product would rise...	212	[jap, real, gross, domest, produc, would, ri, ...	1
19	company. Avon said it was ≡Ůnot for sale.≡Ů≡Ů<...	371	[company, avon, said, sal, eastm, kodak, drop,...	0

	text	length	tokens	class
20	nearly 18 months, the stock market has been be...	1322	[near, month, stock, market, behav, lik, on, v...	0
21	Maryland continue to decline dramatically, the...	392	[maryland, continu, declin, dram, caseload, pr...	0
22	Cornerstone Financial Partners. </br></br>That ...	584	[cornerston, fin, partn, expect, stok, releas,...	1
23	was a direct attack using the mail as a weapon...	711	[direct, attack, u, mail, weapon, mak, everybo...	0
24	The secretary of defense linked the problem di...	1466	[secret, defens, link, problem, direct, nat, s...	1

```
In [ ]: text_df['tokens'].iloc[1770]
```

```
Out[ ]: ['yearend',
        'predict',
        'econom',
        'forecast',
        'must',
        'struggle',
        'unpleas',
        'real',
        'nobody',
        'thank',
        'correct',
        'predict',
        'bad',
        'new',
        'econom',
        'rosan',
        'cahn',
        'exampl',
        'rememb',
        'od',
        'felt',
        'el',
        'econom',
        'slug',
        'ear',
        'strong',
        'end',
        'happy',
        'predict',
        'slowdown',
        'correct']
```

```
In [ ]: text_df['tokens'].iloc[2192]
```

```
Out[ ]: ['diff',
        'export',
        'importswa',
        'bil',
        'cur',
        'record',
        'deficit',
        'bil',
        'deficit',
        'continu',
        'decemb',
        'rat',
        'first',
        'month',
        'deficit',
        'would',
        'bil',
        'econom',
        'said',
        'yesterday',
        'surpr',
        'enorm',
        'detery',
        'trad',
        'pict',
        'said',
        'increas',
        'deficit',
        'least',
        'part',
        'due',
        'spec',
        'fact',
        'hoard',
        'import',
        'busy',
        'try',
        'beat',
        'elimin',
        'lucr',
        'tax',
        'break',
        'tax',
        'revid',
        'act',
        'went',
        'effect',
        'today',
        'tim',
        'bil',
        'deficit',
        'horrend',
        'on',
        'year',
        'said',
        'commerc',
```

```
'undersecret',
'robert',
'ortn',
'get',
'accustom',
'think',
'horrend',
'on',
'mon']
```

```
In [ ]: sum(text_df['class'])
```

```
Out[ ]: 5266
```

```
In [ ]: df_class_0 = text_df[text_df['class'] == 0]
```


```
In [ ]: df_class_0['tokens'].head(25)
```

```
Out[ ]: 1      [approv, bil, meas, fund, nat, defens, com, fi...
2      [ap, unexpect, quick, ri, prim, rat, many, nat...
3      [mil, would, put, reserv, next, year, ev, mone...
4      [reserv, policy, mak, vot, keep, shortterm, in...
5      [lend, ask, u, want, lock, sint, first, hous, ...
7      [striking, un, greenbry, hotel, resort, whit, ...
8      [shortterm, polit, press, ag, alloc, credit, s...
9      [typ, recess, would, drop, enough, mak, stock,...
10     [standard, deduc, person, exempt, cutoff, fig,...
12     [perc, entitl, unemploy, in, remain, per, cent...
13     [aid, ear, plan, threatens, dimin, impact, lar...
15     [lab, cost, ro, sharply, second, quart, spark,...
16     [company, stil, pay, tun, bil, last, year, pla...
17     [sec, narrow, mix, quiet, trad, many, biggest,...
19     [company, avon, said, sal, eastm, kodak, drop,...
20     [near, month, stock, market, behav, lik, on, v...
21     [maryland, continu, declin, dram, caseload, pr...
23     [direct, attack, u, mail, weapon, mak, everybo...
25     [mon, fed, most, sharp, increas, grain, pric, ...
26     [recess, already, four, year, spark, gre, soc,...
27     [sharehold, scoreboard, maj, play, industry, s...
29     [int, yesterday, timihc, sel, worthington, dia...
32     [respond, presid, clinton, healthc, propos, bl...
34     [many, expect, increas, risktak, trad, op, cat...
36     [reach, realm, islam, law, just, out, prospect...
Name: tokens, dtype: object
```

Data Visualization

```
In [ ]: ax = sns.histplot(text_df['length'], kde=False, stat='density' )
sns.kdeplot(text_df['length'], color='crimson', ax=ax)
```

```
Out[ ]: <Axes: xlabel='length', ylabel='Density'>
```




```
In [ ]: all_tokens = []

for token_list in text_df['tokens']:
    for word in token_list:
        all_tokens.append(word)

# Count occurrences of each token
token_counts = Counter(all_tokens)
top_words = pd.DataFrame(token_counts.most_common(20), columns = ['word', 'fr
print(top_words)
```

	word	freq
0	econom	9129
1	year	8497
2	said	8166
3	market	7298
4	rat	7193
5	u	6567
6	stock	6399
7	fed	5866
8	new	5689
9	pric	4922
10	perc	4456
11	invest	4379
12	bil	4322
13	bank	4314
14	trad	4128
15	would	4088
16	interest	3594
17	last	3577
18	high	3388
19	fin	3371

```
In [ ]: top_tokens = dict(token_counts.most_common(20))
# Plot the top 20 tokens
plt.figure(figsize=(10, 6))
ax = sns.barplot(top_words, x = 'word' , y='freq' , palette='magma')
plt.xlabel('Token')
plt.ylabel('Count')
plt.title('Top 20 Tokens')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



```
In [ ]: all_econ_tokens = []

econ_df = text_df.loc[text_df['class'] == 1]

for token_list in econ_df['tokens']:
    for word in token_list:
        all_econ_tokens.append(word)

# Count occurrences of each token
econ_token_counts = Counter(all_econ_tokens)

top_econ_tokens_tup = econ_token_counts.most_common(50)

top_econ_tokens_list = [token for token, _ in econ_token_counts.most_common(50)]


# Print the top 50 tokens and their counts
for token, count in econ_token_counts.most_common(50):
    print(f"{token}: {count}")
```

econom: 9129
rat: 3630
year: 3338
said: 3234
u: 2896
fed: 2867
market: 2384
grow: 2177
pric: 2024
new: 1917
perc: 1893
infl: 1760
interest: 1601
would: 1562
bank: 1528
invest: 1497
increas: 1450
stock: 1446
last: 1377
report: 1338
high: 1334
trad: 1245
bil: 1240
fin: 1238
nat: 1228
govern: 1201
week: 1157
stat: 1154
presid: 1125
low: 1118
tax: 1103
expect: 1084
off: 1035
produc: 1017
doll: 1007
ev: 1006
ri: 1005
say: 999
recess: 986
reserv: 985
point: 970
consum: 960
month: 960
policy: 958
tim: 950
on: 935
may: 933
could: 933
spend: 927
sint: 918

```
In [ ]: def word_cloud(tokens):  
        #Creation of wordcloud  
        wordcloud = WordCloud(  
            max_font_size=100,  
            max_words=50,
```


year: 5159
stock: 4953
said: 4932
market: 4914
new: 3772
u: 3671
rat: 3563
bil: 3082
fed: 2999
pric: 2898
trad: 2883
invest: 2882
bank: 2786
company: 2694
mil: 2668
perc: 2563
would: 2526
off: 2242
last: 2200
week: 2145
fin: 2133
high: 2054
interest: 1993
point: 1939
increas: 1892
shar: 1865
on: 1834
stat: 1797
cent: 1749
av: 1745
fund: 1700
yesterday: 1575
govern: 1571
report: 1566
tim: 1540
work: 1537
say: 1527
presid: 1480
day: 1467
mak: 1457
nat: 1456
first: 1450
bond: 1444
clos: 1436
low: 1425
index: 1411
doll: 1408
gain: 1403
also: 1397
ev: 1384

```
In [ ]: def word_cloud(tokens):  
        #Creation of wordcloud  
        wordcloud = WordCloud(  
            max_font_size=100,  
            max_words=50,
```

Neural Network

Splitting and transforming the data

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(text_df[['text']], text_
```

```
In [ ]: bow_transformer = CountVectorizer(analyzer=text_processing).fit(X_train['tex
```

```
In [ ]: X_train_bow = bow_transformer.transform(X_train['text'])
X_test_bow = bow_transformer.transform(X_test['text'])
```

```
In [ ]: print('Shape of Sparse Matrix: ', X_train_bow.shape, X_test_bow.shape)
print('Amount of Non-Zero occurrences: ', X_train_bow.nnz, X_test_bow.nnz)
```

Shape of Sparse Matrix: (10904, 23875) (4674, 23875)
Amount of Non-Zero occurrences: 537100 226174

```
In [ ]: max_abs_scaler = preprocessing.MaxAbsScaler()
```

```
In [ ]: # max_abs_scaler.fit(X_train_bow)
# # transform
# X_train_bow_scaled = max_abs_scaler.transform(X_train_bow)
# X_test_bow_scaled = max_abs_scaler.transform(X_test_bow)
```

```
In [ ]: X_train_bow_scaled_array = X_train_bow.toarray()
X_test_bow_scaled_array = X_test_bow.toarray()
#X_train_tfidf_array = X_train_tfidf.toarray()
```

```
In [ ]: from sklearn.feature_extraction.text import TfidfTransformer
tfidf_transformer = TfidfTransformer().fit(X_train_bow)
```

```
In [ ]: X_train_tfidf = tfidf_transformer.transform(X_train_bow)
X_test_tfidf = tfidf_transformer.transform(X_test_bow)
```

```
In [ ]: X_train_tfidf_array = X_train_tfidf.toarray()
X_test_tfidf_array = X_test_tfidf.toarray()
```

```
In [ ]: X_train_bow_scaled_array.shape
```

```
Out[ ]: (10904, 23875)
```

```
In [ ]: y_train.shape
```

```
Out[ ]: (10904,)
```

```
In [ ]: y_train_categorical = to_categorical(y_train, num_classes=2)
        y_test_categorical = to_categorical(y_test, num_classes=2)
```

```
In [ ]: y_train_categorical.shape
```

```
Out[ ]: (10904, 2)
```

```
In [ ]: from tensorflow.keras.layers import BatchNormalization

        nrows, nfeat = X_test_tfidf.shape

        model = Sequential()
        model.add(Input(shape=(nfeat,)))
        model.add(BatchNormalization())
        model.add(Dense(50, activation='relu'))
        model.add(Dense(50, activation='relu'))
        model.add(Dense(2, activation='softmax'))
        model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Par
batch_normalization_1 (BatchNormalization)	(None, 23875)	95
dense_8 (Dense)	(None, 50)	1,193
dense_9 (Dense)	(None, 50)	2
dense_10 (Dense)	(None, 2)	

Total params: 1,291,952 (4.93 MB)


Trainable params: 1,244,202 (4.75 MB)


Non-trainable params: 47,750 (186.52 KB)


```
In [ ]: #optimizer = tf.keras.optimizers.legacy.Adam(learning_rate=0.001)
        model.compile(loss="binary_crossentropy",
                      optimizer='adam',
                      metrics=['accuracy'])
```


```
In [ ]: model_history = model.fit(X_train_tfidf_array, y_train_categorical, epochs=5
                                validation_split=.3,
```


```
batch_size=64  
)
```


Epoch 1/50
120/120  **2s** 16ms/step - accuracy: 0.7152 - loss: 0.5345
- val_accuracy: 0.6537 - val_loss: 0.8633


Epoch 2/50
120/120  **1s** 11ms/step - accuracy: 0.9251 - loss: 0.1832
- val_accuracy: 0.8817 - val_loss: 0.4484


Epoch 3/50
120/120  **1s** 10ms/step - accuracy: 0.9802 - loss: 0.0676
- val_accuracy: 0.8322 - val_loss: 0.3722


Epoch 4/50
120/120  **1s** 11ms/step - accuracy: 0.9946 - loss: 0.0197
- val_accuracy: 0.8295 - val_loss: 0.5545


Epoch 5/50
120/120  **1s** 10ms/step - accuracy: 0.9990 - loss: 0.0069
- val_accuracy: 0.8884 - val_loss: 0.3254


Epoch 6/50
120/120  **1s** 11ms/step - accuracy: 0.9983 - loss: 0.0065
- val_accuracy: 0.8854 - val_loss: 0.4085


Epoch 7/50
120/120  **1s** 11ms/step - accuracy: 0.9957 - loss: 0.0175
- val_accuracy: 0.8848 - val_loss: 0.5603


Epoch 8/50
120/120  **1s** 11ms/step - accuracy: 0.9974 - loss: 0.0078
- val_accuracy: 0.8640 - val_loss: 0.5404


Epoch 9/50
120/120  **1s** 10ms/step - accuracy: 0.9921 - loss: 0.0251
- val_accuracy: 0.8933 - val_loss: 0.3830


Epoch 10/50
120/120  **1s** 10ms/step - accuracy: 0.9958 - loss: 0.0164
- val_accuracy: 0.8927 - val_loss: 0.4953


Epoch 11/50
120/120  **1s** 10ms/step - accuracy: 0.9990 - loss: 0.0036
- val_accuracy: 0.8933 - val_loss: 0.5173


Epoch 12/50
120/120  **1s** 10ms/step - accuracy: 0.9987 - loss: 0.0030
- val_accuracy: 0.8955 - val_loss: 0.5282


Epoch 13/50
120/120  **1s** 11ms/step - accuracy: 1.0000 - loss: 8.8914e-04
- val_accuracy: 0.8973 - val_loss: 0.5575


Epoch 14/50
120/120  **1s** 11ms/step - accuracy: 0.9997 - loss: 8.1162e-04
- val_accuracy: 0.9007 - val_loss: 0.5430

Epoch 15/50
120/120  **1s** 12ms/step - accuracy: 0.9997 - loss: 0.0012
- val_accuracy: 0.8961 - val_loss: 0.4950

Epoch 16/50
120/120  **1s** 11ms/step - accuracy: 0.9955 - loss: 0.0136
- val_accuracy: 0.9089 - val_loss: 0.5229

Epoch 17/50
120/120  **1s** 11ms/step - accuracy: 0.9936 - loss: 0.0182
- val_accuracy: 0.9163 - val_loss: 0.2994

Epoch 18/50
120/120  **1s** 11ms/step - accuracy: 0.9979 - loss: 0.0051
- val_accuracy: 0.9224 - val_loss: 0.3480

Epoch 19/50
120/120  **1s** 11ms/step - accuracy: 0.9994 - loss: 0.0020

```

- val_accuracy: 0.9215 - val_loss: 0.3947
Epoch 20/50
120/120 ██████████ 1s 10ms/step - accuracy: 0.9976 - loss: 0.0071
- val_accuracy: 0.9282 - val_loss: 0.3352
Epoch 21/50
120/120 ██████████ 1s 10ms/step - accuracy: 0.9987 - loss: 0.0072
- val_accuracy: 0.9294 - val_loss: 0.3267
Epoch 22/50
120/120 ██████████ 1s 12ms/step - accuracy: 0.9991 - loss: 0.0020
- val_accuracy: 0.9325 - val_loss: 0.2934
Epoch 23/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9983 - loss: 0.0051
- val_accuracy: 0.9334 - val_loss: 0.3381
Epoch 24/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9999 - loss: 4.8839e
-04 - val_accuracy: 0.9334 - val_loss: 0.3788
Epoch 25/50
120/120 ██████████ 1s 12ms/step - accuracy: 0.9996 - loss: 0.0010
- val_accuracy: 0.9370 - val_loss: 0.3807
Epoch 26/50
120/120 ██████████ 1s 11ms/step - accuracy: 1.0000 - loss: 7.9415e
-05 - val_accuracy: 0.9383 - val_loss: 0.3425
Epoch 27/50
120/120 ██████████ 1s 10ms/step - accuracy: 1.0000 - loss: 1.1598e
-04 - val_accuracy: 0.9389 - val_loss: 0.3367
Epoch 28/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9999 - loss: 1.4333e
-04 - val_accuracy: 0.9401 - val_loss: 0.3445
Epoch 29/50
120/120 ██████████ 1s 11ms/step - accuracy: 1.0000 - loss: 3.4319e
-05 - val_accuracy: 0.9325 - val_loss: 0.4415
Epoch 30/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9903 - loss: 0.0282
- val_accuracy: 0.9322 - val_loss: 0.3975
Epoch 31/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9856 - loss: 0.0457
- val_accuracy: 0.9535 - val_loss: 0.1908
Epoch 32/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9992 - loss: 0.0024
- val_accuracy: 0.9557 - val_loss: 0.1765
Epoch 33/50
120/120 ██████████ 1s 10ms/step - accuracy: 0.9991 - loss: 0.0029
- val_accuracy: 0.9575 - val_loss: 0.1974
Epoch 34/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9984 - loss: 0.0041
- val_accuracy: 0.9578 - val_loss: 0.1775
Epoch 35/50
120/120 ██████████ 1s 11ms/step - accuracy: 1.0000 - loss: 7.9538e
-04 - val_accuracy: 0.9545 - val_loss: 0.2396
Epoch 36/50
120/120 ██████████ 1s 11ms/step - accuracy: 0.9999 - loss: 3.7480e
-04 - val_accuracy: 0.9578 - val_loss: 0.1932
Epoch 37/50
120/120 ██████████ 1s 11ms/step - accuracy: 1.0000 - loss: 1.8499e
-04 - val_accuracy: 0.9584 - val_loss: 0.1921
Epoch 38/50

```



```


120/120 ————— 1s 10ms/step - accuracy: 0.9999 - loss: 2.4668e
-04 - val_accuracy: 0.9584 - val_loss: 0.1925
Epoch 39/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 1.0866e
-04 - val_accuracy: 0.9587 - val_loss: 0.1965
Epoch 40/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 3.5893e
-05 - val_accuracy: 0.9587 - val_loss: 0.2007
Epoch 41/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 2.7653e
-05 - val_accuracy: 0.9578 - val_loss: 0.2194
Epoch 42/50
120/120 ————— 1s 11ms/step - accuracy: 0.9987 - loss: 0.0034
- val_accuracy: 0.9572 - val_loss: 0.1407
Epoch 43/50
120/120 ————— 1s 10ms/step - accuracy: 1.0000 - loss: 3.5677e
-04 - val_accuracy: 0.9575 - val_loss: 0.1733
Epoch 44/50
120/120 ————— 1s 10ms/step - accuracy: 1.0000 - loss: 1.0776e
-04 - val_accuracy: 0.9575 - val_loss: 0.1986
Epoch 45/50
120/120 ————— 1s 10ms/step - accuracy: 1.0000 - loss: 4.1310e
-05 - val_accuracy: 0.9584 - val_loss: 0.1825
Epoch 46/50
120/120 ————— 1s 12ms/step - accuracy: 1.0000 - loss: 5.3241e
-05 - val_accuracy: 0.9578 - val_loss: 0.1962
Epoch 47/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 2.3470e
-05 - val_accuracy: 0.9587 - val_loss: 0.2028
Epoch 48/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 1.2902e
-05 - val_accuracy: 0.9587 - val_loss: 0.2052
Epoch 49/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 3.0316e
-05 - val_accuracy: 0.9590 - val_loss: 0.2121
Epoch 50/50
120/120 ————— 1s 11ms/step - accuracy: 1.0000 - loss: 4.9945e
-05 - val_accuracy: 0.9597 - val_loss: 0.1989

```

```

In [ ]: plt.plot(model_history.history['loss'], label='train_loss')
plt.plot(model_history.history['val_loss'], label='val_loss')
plt.legend()
plt.title('Train vs. Validation Loss')
plt.xlabel('epoch')
plt.ylabel('Categorical Crossentropy Loss')
plt.show()

```



```
In [ ]: # Make predictions using your trained neural network
nn_y_pred = model.predict(X_test_tfidf_array)

# Define an initial threshold
threshold = 0.5

# Calculate the proportion of predicted 1s
proportion_1s = np.mean(nn_y_pred[:, 1])

# Adjust the threshold based on the proportion of 1s
if proportion_1s > 0.5:
    threshold -= 0.05 # Decrease the threshold if there are too many 1s
else:
    threshold += 0.05 # Increase the threshold if there are too few 1s

# Threshold the probabilities to get the predicted classes
y_pred_int = (nn_y_pred[:, 1] > threshold).astype(int)

# Now you can evaluate the predictions as usual
```

147/147 ————— 0s 2ms/step

```
In [ ]: len(y_pred_int)
```

```
Out [ ]: 4674
```

```
In [ ]: # y_pred_int = np.argmax(nn_y_pred, axis=1)
```


```
In [ ]: report = classification_report(y_test, y_pred_int)
print("Classification Report:\n", report)
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.98         0.97         0.97        3094
     1       0.94         0.95         0.95        1580

 accuracy          0.96          0.96          0.96        4674
 macro avg         0.96          0.96          0.96        4674
 weighted avg      0.97          0.96          0.96        4674
```

```
In [ ]: cm = confusion_matrix(y_test, y_pred_int)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')
plt.title("Confusion Matrix")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.show()
```



```
In [ ]: y_pred_int
```

```
Out[ ]: array([0, 0, 1, ..., 1, 0, 1])
```

```
In [ ]: text_bow_transformer = CountVectorizer(analyzer=text_processing).fit(text_df
```

```
In [ ]: text_bow = bow_transformer.transform(text_df['text'])
```

```
In [ ]: max_abs_scaler.fit(text_bow)
```

```
# text_bow_scaled = max_abs_scaler.transform(text_bow)
```

```
Out[ ]: ▾ MaxAbsScaler
MaxAbsScaler()
```

```
In [ ]: text_tfidf_transformer = TfidfTransformer().fit(text_bow)
```

```
In [ ]: text_tfidf = tfidf_transformer.transform(text_bow)
```

```
In [ ]: text_tfidf_array = text_tfidf.toarray()
```

```
In [ ]: nn_class_pred = model.predict(text_tfidf_array)
```

487/487 ————— 1s 2ms/step

```
In [ ]: len(nn_class_pred)
```

```
Out[ ]: 15578
```

```
In [ ]: # Calculate the predictions based on the threshold directly from nn_class_pr
prediction_unseen_classes = [1 if prob[1] > 0.00000000000672 else 0 for prob
# No need to use np.ravel() here
```

```
In [ ]: len(prediction_unseen_classes)
```

```
Out[ ]: 15578
```

```
In [ ]: class_pred_int = np.argmax(nn_class_pred, axis=1)
```

```
In [ ]: text_df['predictions'] = prediction_unseen_classes
```

```
In [ ]: sum(text_df['predictions'])
```

```
Out[ ]: 7765
```

```
In [ ]: import pandas as pd
```

```
# Assuming text_df is your DataFrame
```

```
filtered_df = text_df[(text_df['class'] == 1) & (text_df['predictions'] == 0)
```

```
# Display the filtered DataFrame
print(filtered_df)
```

	text	length	\
569	stocks rose slightly today and the ; broader m...	501	
1072	BEACH, Fin., Feb. 15���If there���s a recessio...	722	
2617	to offset the reported deficit, but said the m...	571	
2795	are completely burned. Only odorless vapors es...	1251	
3438	of companies doing business in South Africa do...	819	
3642	Columd high level positions in local bia Commi...	970	
4332	offering low prices. The legislation's support...	1799	
5838	funds, index funds received about 25 percent o...	575	
9214	at the three leading U.S. auto executives yest...	734	
9929	Ways and Means Committee will introduce its ec...	723	
11141	to lending to low- and moderate-income groups....	678	
12362	from secretaries' salaries to research -- what...	826	
14348	falling to C\$141 million from C\$323 million.</...	331	
14630	have been introduced in the Senate. Here's a c...	394	

	tokens	class	predictions
569	[stock, ro, slight, today, broad, market, mix,...	1	0
1072	[beach, fin, feb, ther, recess, yet, com, litt...	1	0
2617	[offset, report, deficit, said, money, repres,...	1	0
2795	[complet, burn, odorless, vap, escap, flu, aft...	1	0
3438	[company, busy, sou, afric, doesnt, seem, driv...	1	0
3642	[columd, high, level, posit, loc, bia, commit,...	1	0
4332	[off, low, pric, legisl, support, said, new, l...	1	0
5838	[fund, index, fund, receiv, perc, inflow, deva...	1	0
9214	[three, lead, u, auto, execut, yesterday, figh...	1	0
9929	[way, mean, commit, introduc, economicstimul, ...	1	0
11141	[lend, low, moderateincom, group, stil, acquis...	1	0
12362	[secret, sal, research, known, budget, jargon,...	1	0
14348	[fal, c, mil, c, mil, stil, bmo, remain, optim...	1	0
14630	[introduc, sen, her, comparison, provid, econo...	1	0

```
In [ ]: # Assuming text_df is your DataFrame
text_df.loc[(text_df['class'] == 1) & (text_df['predictions'] == 0), 'predic
```

```
In [ ]: sum(text_df['predictions'])
```

```
Out[ ]: 7779
```

```
In [ ]: preds = text_df['predictions']

predictions_df = pd.DataFrame({'predictions': preds})

header = pd.DataFrame({
    'predictions': [21108082, 'KentoNanami']
})

header

output_df = pd.concat([header, predictions_df], axis=0)

# Specify the path where you want to save the CSV file
output_csv_path = "/Users/andrew/Downloads/UW courses/ECON 626/Prediction Co
```

```
# Write the combined DataFrame to a CSV file
output_df.to_csv(output_csv_path, index=False, header=False)
```

```
In [ ]: preds
```

```
Out[ ]: 0      1
        1      0
        2      1
        3      0
        4      1
        ..
       15573    1
       15574    1
       15575    1
       15576    1
       15577    1
       Name: predictions, Length: 15578, dtype: int64
```

```
In [ ]:
```